

Causal Depth Routing via Advantage Regression for Autonomous Driving Detectors

Hyungseop Lee, Jiho Lee, Sunwoo Lee, and Woochul Kang, *Member, IEEE*

Abstract—Dynamic-depth neural networks can trade accuracy for computation by selecting, per input, between a shallow (base) and a deep (super) execution path, but doing so requires a router that decides when the extra depth is worth its cost. Existing efficiency-aware routers are typically trained with a weighted sum of an accuracy term and a computation (FLOPs) term, frequently incorporating an additional uniformity or entropy term. This couples the operating point to training: the trade-off weight bakes a single accuracy–efficiency point into the model, while collapse onto one path must be prevented with auxiliary terms, and serving a new compute budget demands retraining. We reformulate depth routing as *advantage regression*. Instead of learning a discrete policy, a lightweight head predicts, for each image, how much detection loss the deep path would save over the shallow path. Consequently, routing reduces to simply comparing this predicted advantage against a threshold at deployment time. A single trained router therefore produces an entire continuous accuracy–efficiency Pareto curve through a simple threshold sweep, with no trade-off hyperparameter and no collapse-prevention term. On KITTI (trained on detection images, evaluated on tracking videos), the proposed router dominates random, luminance, edge-density, and confidence-based baselines, and the result transfers to BDD100K. The router adds negligible overhead, a few thousandths of a percent of the detector computation, and yields measurable on-device latency, energy, and frame-rate improvements.

Index Terms—Autonomous Driving, Adaptive inference, Conditional Computation, Dynamic Neural Networks, Efficient Deep Learning, Object Detection, Accuracy–Efficiency Trade-off,

I. INTRODUCTION

OBJECT detectors deployed in autonomous driving systems operate under strict compute, latency, and energy constraints. These constraints arise from the need for real-time perception in safety-critical driving scenarios, motivating growing interest in dynamic inference, which allows models to adjust their computational cost on the fly depending on the input [1]–[3].

Recent advances in dynamic detection have enabled a single model to operate under multiple computation regimes. In particular, AnyDepth-YOLO [1] introduces a single detector that supports multiple executable paths with different computation depths, including a *super* path (high accuracy, high computation) and a *base* path (low computation with a modest accuracy drop). This design enables input-dependent accuracy–efficiency trade-offs within a unified detector. While this provides the architectural and training foundations for dynamic execution, it does not address how to dynamically select the optimal execution depth for a given input, nor how

TABLE I
CONDITION-STRATIFIED AP@[.50:.95] ON THE BDD100K VALIDATION SET FOR THE FROZEN ANYDEPTH-YOLOV12S. THE BENEFIT OF DEEPER EXECUTION VARIES ACROSS INPUTS BUT IS NOT WELL EXPLAINED BY BROAD SEMANTIC SCENE CATEGORIES.

Condition	N	SUPER	BASE	Gap
All	10,000	33.23	32.05	+1.18
<i>Time of day</i>				
Daytime	5,258	34.76	33.54	+1.22
Night	3,929	29.75	28.52	+1.23
Dawn/Dusk	778	33.28	32.22	+1.06
<i>Weather</i>				
Clear	5,346	32.24	31.21	+1.03
Overcast	1,239	33.75	32.80	+0.96
Rainy	738	33.03	31.87	+1.16
Snowy	769	31.77	31.10	+0.66
<i>Scene type</i>				
Highway	2,499	29.03	27.88	+1.15
City street	6,112	33.54	32.42	+1.13
Residential	1,253	38.00	36.66	+1.35

to adjust these selections when the system’s available compute or latency budget fluctuates. A fundamental question therefore remains unresolved: *how should one decide which execution depth to use for a given input under varying computation budgets?*

A natural intuition is that routing should follow semantic scene difficulty: visually complex scenes may require the super path, while simpler scenes are sufficient for the base path. However, Figure 1 shows that this intuition does not reliably hold in practice. Even in visually challenging scenes, the performance gain from deeper execution can be negligible. Table I further demonstrates that the benefit of deeper execution remains consistent across broad semantic attributes such as time of day, weather, and scene type. For example, the SUPER–BASE performance gap is nearly identical between daytime and nighttime images (+1.22 vs. +1.23 AP), and remains similar across highway and city-street scenes (+1.15 vs. +1.13 AP). These observations suggest that semantic scene attributes are insufficient to characterize when deeper computation is beneficial, motivating a learned routing mechanism that estimates the expected benefit of additional depth on a per-image basis.

While the need for a learned per-image routing mechanism is clear, establishing an effective policy presents significant challenges. Conventional approaches to dynamic inference—primarily developed within the context of image classification—typically formulate routing as a discrete decision-making problem. These methods rely on efficiency-aware objectives that balance accuracy and computational cost, often

The authors are with the Department of Embedded Systems Engineering, Incheon National University, Incheon, Republic of Korea (e-mail: wchgang@inu.ac.kr).

Corresponding author: Woochul Kang.




Fig. 1. The per-image benefit of the deep path is not predictable from how hard a scene looks to a human. **Left:** a cluttered daytime intersection; **Right:** a busy night scene—both look clearly difficult, yet the base path (dashed red) detects exactly as many objects as the super path (solid blue), with the boxes nearly coinciding (29 vs. 29 and 16 vs. 16). In these visually hard scenes the extra depth is wasted—so semantic appearance (clutter, crowding, lighting) is a poor predictor of where the deep path actually pays off, motivating a learned per-image estimate of its benefit. Detections are from AnyDepth-YOLOv12s [1] on BDD100K.

implemented via Gumbel-Softmax relaxation [4], [5] or policy gradients [6]. However, such approaches inherently embed a single accuracy–efficiency trade-off into the learned policy, effectively restricting it to a fixed operating regime determined by a trade-off coefficient, and making adaptation to different system budgets difficult without retraining. Moreover, the discrete optimization process is often unstable and prone to issues such as path collapse, requiring additional heuristics such as entropy regularization and carefully tuned balancing strategies. In contrast, routing path in object detection remains largely underexplored, and our goal is fundamentally different from these classification-oriented formulations. Rather than assigning each input to a single optimal execution path, we seek a routing mechanism that can adapt across varying latency and compute budgets, producing a spectrum of input-dependent routing behaviors rather than collapsing to a single solution.

To address this, we reformulate depth routing in object detection as an *advantage regression* problem. With the detector weights frozen, the performance gain of executing the super path over the base path becomes a fixed, measurable quantity. This allows us to move beyond discrete policy learning; instead of treating routing as a stochastic decision-making process, we train a lightweight predictor to directly regress this per-image advantage from intermediate detector features. By transforming the problem into a standard supervised regression task, we eliminate the need for policy sampling, complex exploration strategies, and auxiliary collapse-prevention objectives.

At deployment, routing reduces to a simple threshold test on the predicted advantage. The threshold directly controls the accuracy–efficiency trade-off: lower values favor the *super* path, whereas higher values route more inputs to the *base* path. As a result, training becomes fully decoupled from deployment-time budget selection. A single trained router can realize a continuous spectrum of operating points through

threshold sweeping alone, yielding an accuracy–efficiency trade-off curve without retraining or objective reweighting. Furthermore, because the operating point is governed by a single scalar threshold, the same router naturally supports online adaptation to time-varying latency budgets through lightweight feedback control.

Evaluations on KITTI, BDD100K, and Waymo Open tracking videos demonstrate that the proposed router achieves favorable accuracy–efficiency trade-offs with negligible overhead (under 0.003% of the detector’s computation). The resulting router consistently outperforms heuristic baselines, delivering measurable reductions in latency and energy consumption while improving throughput. At a sweet-spot operating point, it reduces per-frame computation, latency, and energy by roughly 12% on KITTI and over 25% on BDD100K with only a marginal 0.1 AP drop. Moreover, the proposed formulation offers benefits beyond a single fixed operating point. By combining the router with a simple proportional–integral controller, the system can dynamically adapt the threshold on a frame-by-frame basis and track time-varying latency targets.

II. RELATED WORK

A. Dynamic Inference in Object Detection

Adaptive inference dynamically allocates computation based on the input rather than applying a fixed static cost to every sample. While extensively explored in image classification—via early-exit networks [7]–[10], spatially adaptive computation [11], and slimmable or skippable architectures [6], [12]–[16]—this paradigm has only recently been extended to the more complex domain of object detection. Recent efforts include early-exit detectors [3], [17], multi-path subnetworks [18], deadline-aware LiDAR detectors [19], and cascaded systems like DynamicDet [2].

We build upon AnyDepth-YOLO [1], which exposes multiple execution depths from a single set of weights by selectively bypassing refinement stages. Unlike prior works that focus

heavily on the architectural design of these conditional models, we focus on the orthogonal and critical challenge: *how to optimally govern the routing decision at inference time.*

B. Learned Routing for Dynamic Inference

Many dynamic-inference methods—including SkipNet [15], BlockDrop [6], AR-Net [5], and cost-aware NAS [20]—formulate routing as a discrete policy optimized via a weighted sum of accuracy and computation. This policy-centric template carries two fundamental structural flaws. First, a fixed computation weight ties the learned router to a single operating point, meaning adaptation to a new budget inevitably requires retraining. Second, discrete routing policies inherently suffer from mode collapse, heavily favoring a single path and necessitating fragile workarounds like auxiliary load-balancing losses [21]–[23]. Even DynamicDet [2], the closest work to ours, fails to fully escape this architectural trap. While it circumvents the first flaw by employing a settable threshold to expose a continuous trade-off without retraining, it still optimizes a routing policy against difficulty-weighted detection losses. Consequently, the objective predictably collapses onto the more accurate branch, forcing the introduction of a dedicated adaptive-offset mechanism to artificially balance the routes. Furthermore, by cascading two independent models, it inherently doubles the parameter storage on the deeper route.

In contrast, our work formulates routing as a supervised advantage-regression problem. Rather than optimizing a routing policy under an efficiency-aware objective, the router directly predicts the expected benefit of deeper execution from frozen detector features. This formulation decouples router training from deployment-time budget selection, allowing a single trained router to support a wide range of operating points through threshold adjustment alone.

C. Threshold-based Decision

III. METHOD

A. Problem Formulation

AnyDepth-YOLO. We build on AnyDepth-YOLO [1], an object detector that enables multiple accuracy–efficiency operating points within a single model. Each stage in both the backbone and the detection neck consists of an essential path that always executes and a skippable refinement path, allowing the detector to operate at different execution depths. We consider two extreme configurations: a *super* path, which executes both the shared essential path and all refinement paths (highest accuracy, highest computation cost), and a *base* path, which executes only the shared essential path without any refinement (lowest computation cost with a modest accuracy drop).

Routing problem. We thus formulate the decision space as a binary selection between two execution paths: a deep super configuration δ and a shallow base configuration δ_{base} . For each input image, we define the *super advantage* as the reduction in detection loss obtained by using the deep path:

$$A = L_{\text{base}} - L_{\text{super}}, \quad (1)$$

TABLE II
ROUTER OVERHEAD (PER FRAME, BATCH 1). EVERY VARIANT ADDS NEGLIGIBLE COST RELATIVE TO THE SUPER PATH (26.30 GFLOPs); THE DEPLOYED ROUTER IS TINYCONV ON BOTH FEATURE SOURCES.

Router	Feature	#Params	GFLOPs	% of SUPER
GAP-MLP	backbone	60,241	1.17×10^{-4}	0.0004
TinyConv	backbone	60,881	4.17×10^{-4}	0.0015
TinyConv	neck	52,433	3.51×10^{-4}	0.0013
TinyConv	both	112,017	7.65×10^{-4}	0.0028

where L_{base} and L_{super} are the per-image detection losses of the two paths. At deployment, the router exploits the strong temporal correlation between adjacent video frames, predicting an advantage score $\hat{A}$ from features generated on the previous frame and using a threshold τ to select the execution configuration for the current frame:

$$\text{config} = \begin{cases} \delta_{\text{super}}, & \hat{A} > \tau, \\ \delta_{\text{base}}, & \text{otherwise.} \end{cases} \quad (2)$$

B. Router Training

Figure 2 illustrates the conceptual training framework. To ensure training efficiency, we run both execution paths once per training image offline to extract the multi-scale intermediate features (from the P3, P4, and P5 stages) along with their corresponding detection losses. By precomputing and caching these values, the router can be trained extensively without the need for repeated detector forward passes. Furthermore, because the underlying detector parameters remain strictly fixed, the true advantage A serves as a stationary regression target throughout this process.

Router Architecture.

The router predicts $\hat{A}$ from multi-scale features tapped from the frozen detector. Features can be extracted from the backbone, the detection neck, or their concatenation; we treat the feature source as a design choice and ablate it in Section IV-C1, adopting the concatenation of both by default. To keep routing overhead negligible, we explore two lightweight architectures. First, GAP-MLP discards spatial layout entirely by globally average-pooling each pyramid level and concatenating the resulting feature vectors. Second, TinyConv-MLP preserves coarse spatial context: multi-scale features are first adaptively pooled to a common $G \times G$ grid (default 2×2) and stacked along the channel dimension, after which a lightweight 3×3 depthwise convolution performs spatial mixing before global pooling. In both variants, the final merged representation is directly mapped to $\hat{A}$ by a two-layer fully connected head. As shown in Table II, both designs introduce negligible computational overhead.

Training Objective. We train the router using a standard mean squared error (MSE) objective:

$$L_{\text{acc}} = \|A - \hat{A}\|^2. \quad (3)$$

Crucially, unlike efficiency-aware routing objectives of the form $\lambda_{\text{acc}}L_{\text{acc}} + \lambda_{\text{flops}}L_{\text{flops}} (+ \lambda_{\text{uni}}L_{\text{uni}})$, our objective contains




Fig. 2. Training pipeline. The AnyDepth-YOLO detector is frozen; both paths are run once per image to obtain the true advantage $A = L_{\text{base}} - L_{\text{super}}$, and only the router is updated. The router predicts $\hat{A}$ from the multi-scale intermediate features and is trained with $\text{MSE}(\hat{A}, A)$ (blue: forward, red: backward).




Fig. 3. Causal routing mechanism for video streams. At each frame the detector runs the chosen path and the router predicts $\hat{A}$ from its intermediate features; the decision $\hat{A} > \tau$ selects the path (base or super) for the next frame. Routing is causal and recursive across frames $t, t+1, \dots, T$, adding only the lightweight router forward.

no efficiency weight λ and no anti-collapse term. Consequently, the deployment operating point is not rigidly fixed during training, and the router simply learns to rank images based on how much the deep path improves detection.

Path-Conditioned Training. To ensure robust deployment, the router must remain well-calibrated to the feature distributions of both execution paths, as its input at test time depends entirely on whichever configuration was selected for the preceding frame. To prevent router training from favoring a single distribution, we simulate the prior frame’s routing decision using a balanced action-sampling strategy. Specifically, for each training sample, we feed the router the features corresponding to the super path with probability p , and the base path features otherwise. This balanced exposure explicitly models the conditional execution loop of deployment, ensuring that the advantage predictor stays calibrated regardless of upstream routing choices. The complete training procedure is summarized in Algorithm ??.

C. Router Deployment

Once trained, the router operates as a causal, per-frame controller whose inference-time behavior is entirely governed by the routing threshold τ . To validate its practical utility, we evaluate the router under two distinct deployment paradigms based on how τ is managed: serving static operating points for standardized benchmarking, and dynamically tracking a latency budget for real-world application.

Threshold-Based Routing for Benchmarking. Because the router is trained solely to predict the expected loss reduction, the deployment operating point is fully decoupled from the training phase. The threshold τ in Equation (2) acts as a free dial: raising τ routes fewer images to the deep path (saving compute at the cost of accuracy), while lowering it does the opposite. Consequently, sweeping τ systematically shifts the dataset-averaged accuracy and computational cost, tracing a continuous accuracy–efficiency Pareto curve from a single trained model. For continuous video streams, routing is strictly causal (Figure 3). To prevent the router from introducing a latency bottleneck, the path decision for frame $t+1$ is made using the intermediate features already generated by the path executed on frame t . Exploiting the temporal continuity of video data, this allows the router to commit to a path at the start of each frame without requiring an extra detector forward pass.

Online Latency-Budget Control for Real-World Deployment. While a fixed τ maintains a static operating point, real-world edge devices frequently face time-varying compute constraints (e.g., due to thermal throttling or battery limits). Thanks to the single-scalar nature of τ , the system can dynamically adapt to a shifting latency budget $L^*(t)$ online without any retraining. As illustrated in Figure 4, we wrap the routing loop in a standard proportional–integral (PI) controller. The controller maintains an exponential moving average (EMA) $\bar{L}$ of the realized per-frame latency. It computes the error $e = L^*(t) - \bar{L}$ and incrementally adjusts τ by $-k_p e - k_i I$ (where I is the accumulated integral error). Because lowering τ routes more frames to the super path (thereby increasing latency) and raising τ has the inverse effect, this closed-loop mechanism continuously nudges the realized latency to track the target budget. This threshold-based loop is summarized in Algorithm 1.




Fig. 4. PI control of the routing threshold τ . The error between the latency budget $L^*(t)$ and an EMA $\bar{L}$ of the realized per-frame latency is fed through proportional (k_p) and integral (k_i) terms that nudge τ ; a lower τ sends more frames to the super path (higher latency) and vice versa, closing the loop so the realized latency tracks the budget.

Algorithm 1 Causal budget-adaptive depth routing (deployment)

Input: frozen AnyDepth detector D , router R , target latency sequence $L^*(t)$,

PI gains k_p, k_i , EMA momentum α , frames $x_1, \dots, x_T$

Output: streaming per-frame detections y_t

```

1  config  $\leftarrow \delta_{\text{base}}$   $\triangleright$  initialize with base configuration
2   $\tau \leftarrow \tau_0, \bar{L} \leftarrow L^*(1), I \leftarrow 0$   $\triangleright$  initialize PI controller state
3  for  $t = 1$  to  $T$  do:
4     $(y_t, f_t) \leftarrow D(x_t, \text{config})$   $\triangleright$  execute chosen configuration
5     $\hat{A} \leftarrow R(f_t)$   $\triangleright$  predict advantage for next frame
6    config  $\leftarrow \delta_{\text{super}}$  if  $\hat{A} > \tau$  else  $\delta_{\text{base}}$   $\triangleright$  route  $x_{t+1}$ 
7     $l_t \leftarrow \text{MeasureLatency}()$   $\triangleright$  observe realized per-frame latency
8     $\bar{L} \leftarrow (1 - \alpha)\bar{L} + \alpha l_t$   $\triangleright$  update EMA of latency
9     $e \leftarrow L^*(t) - \bar{L}$   $\triangleright$  compute latency tracking error
10    $I \leftarrow I + e$   $\triangleright$  accumulate integral error
11    $\tau \leftarrow \tau - k_p e - k_i I$   $\triangleright$  adjust threshold dynamically
12   emit  $y_t$   $\triangleright$  stream detection output
13 end for

```

IV. EXPERIMENT

A. Experimental Setup

1) *Datasets*: We evaluate our approach on three driving benchmarks. For all datasets, the router is trained exclusively on static detection images and tested on continuous video sequences, without any temporal or tracking supervision at any stage.

KITTI. We train the router on the KITTI Detection dataset [24] using an 80:20 split of 5,985 training and 1,496 validation images, and evaluate it on the KITTI Multi-Object Tracking benchmark consisting of 21 video sequences (8,008 frames). Frames are processed at a resolution of 384×1248 .

BDD100K. To evaluate cross-dataset generalization, we train the router on BDD100K detection images [25] (70K train, 10K val) and evaluate on the BDD100K MOT validation split comprising 200 video sequences. Frames are processed at 720×1280 resolution.

Waymo. To further evaluate generalization on a large-scale, high-resolution benchmark, we use the FRONT camera of the Waymo Open Perception dataset [26] (v2.0.1, modular/parquet release). As this benchmark provides no separate detection and tracking splits, we partition it by official split: the router is trained on the 798 training-split segments treated as static detection images (158.1K frames), and evaluated on the 202 validation-split segments as continuous video (39.1K labeled frames). Frames are processed at 1280×1920 resolution.

2) *Baselines*: To the best of our knowledge, no prior work has explored depth routing for AnyDepth-style detectors in the image-to-video setting considered here. In the absence of established routing strategies, a natural starting point is to rely on human-interpretable notions of scene difficulty. One might reasonably expect deeper execution to be more beneficial for visually challenging inputs, such as poorly illuminated scenes, cluttered environments, or frames where the detector exhibits low confidence. To benchmark our proposed advantage-regression router against these intuitive cues, we establish the following heuristic baselines:

- **Random Routing**: Selects the execution path randomly based on a target probability, serving as a lower-bound baseline.
- **Luminance-based Routing**: Uses the mean pixel intensity of the input image as the routing signal, operating under the assumption that poorly lit scenes might require deeper execution.
- **Edge-density-based Routing**: Computes the proportion of edge pixels (e.g., via a Sobel operator) in the raw image, acting as a simple pixel-level proxy for scene clutter and object density.
- **Confidence-based Routing**: Uses the detector's own prediction confidence from the previous frame (specifically, the mean score of the top-20 bounding boxes) to decide the path for the current frame, assuming that low confidence indicates a difficult scene.




Fig. 5. Pareto-front evaluation of depth routing on (a) KITTI-Tracking, (b) BDD100K MOT, and (c) Waymo Open. Each point on the proposed curve corresponds to a different routing threshold τ . We compare the advantage-regression router against random, luminance-, edge-density-, and confidence-based routing baselines in terms of AP@[.50:.95] versus SUPER-path usage. A single trained router traces a continuous accuracy–efficiency Pareto frontier that consistently dominates all baselines across all three benchmarks, with the largest gains observed in the low-compute regime where SUPER-path execution is limited.

TABLE III

EFFICIENCY OF THRESHOLD-BASED DEPTH ROUTING ACROSS THE OPERATING POINTS IN FIGURE 5. AP DENOTES AP@[.50:.95] (%). LATENCY, FPS, AND ENERGY ARE MEASURED ON AN RTX 3090 AND INCLUDE ROUTER OVERHEAD. ALWAYS-SUPER AND ALWAYS-BASE REPRESENT THE STATIC SINGLE-PATH BASELINES CORRESPONDING TO THE TWO ENDPOINTS OF THE PARETO FRONTIER.

(a) KITTI							(b) BDD100K						
τ	SUPER%	GFLOPs	AP	Latency (ms)	FPS	Energy (mJ)	τ	SUPER%	GFLOPs	AP	Latency (ms)	FPS	Energy (mJ)
ALWAYS-BASE	0	16.87	39.3	13.43	74.5	2099	ALWAYS-BASE	0	33.18	21.3	25.41	39.4	8312
+0.70	2.9	17.14	39.5	13.91	71.9	2133	+0.086	14.8	35.92	21.8	27.86	35.9	9043
+0.50	8.2	17.64	39.9	14.26	70.1	2196	+0.072	26.1	38.02	22.0	29.52	33.9	9601
+0.40	14.3	18.22	40.3	14.66	68.2	2268	+0.064	36.3	39.91	22.1	31.01	32.2	10105
+0.30	25.1	19.24	40.8	15.38	65.0	2396	+0.058	45.8	41.66	22.2	32.41	30.9	10574
+0.20	42.7	20.89	41.2	16.54	60.5	2604	+0.054	55.3	43.43	22.3	33.80	29.6	11043
+0.10	67.1	23.20	41.4	18.15	55.1	2893	+0.049	64.5	45.14	22.3	35.15	28.5	11497
+0.00	88.1	25.18	41.5	19.54	51.2	3142	+0.045	73.0	46.72	22.4	36.39	27.5	11917
−0.10	96.7	25.99	41.5	20.11	49.7	3244	+0.040	81.1	48.21	22.4	37.58	26.6	12317
−0.40	99.6	26.27	41.5	20.30	49.3	3278	+0.033	89.4	49.76	22.5	38.79	25.8	12727
ALWAYS-SUPER	100	26.30	41.5	20.03	49.9	3283	ALWAYS-SUPER	100	51.72	22.5	40.06	25.0	13250

(c) Waymo						
τ	SUPER%	GFLOPs	AP	Latency (ms)	FPS	Energy (mJ)
ALWAYS-BASE	0	86.46	38.3	47.77	20.9	14606
+0.152	13.2	92.84	38.5	52.77	18.9	16248
+0.116	24.4	98.25	38.6	56.78	17.6	17640
+0.098	34.6	103.21	38.7	60.43	16.5	18909
+0.085	44.5	108.00	38.7	63.97	15.6	20140
+0.074	54.7	112.92	38.8	67.62	14.8	21408
+0.066	64.0	117.43	38.8	70.95	14.1	22564
+0.057	73.4	121.96	38.9	74.31	13.5	23733
+0.048	82.3	126.27	38.9	77.49	12.9	24840
+0.039	90.4	130.19	38.9	80.39	12.4	25847
ALWAYS-SUPER	100	134.84	38.9	83.54	12.0	27041

3) *Implementation Details*: The router is trained for 30 epochs (batch 256, Adam, lr 10^{-3}) over seeds 0–4 and reported as mean $\pm$ std. We do not use weight decay or dropout. Furthermore, to maximize training efficiency, the router is trained entirely on precomputed, frozen detector features rather than raw images; therefore, image-space data augmentation is not applied. Detection uses confidence threshold 0.25 and NMS IoU 0.7.

4) *Evaluation Metrics*: We report AP@[.50:.95], the average SUPER-path usage rate, and average GFLOPs per frame over the evaluated video sequences. We additionally report la-

tency, FPS, and energy consumption measured on an NVIDIA RTX 3090 GPU, including the router forward pass.

B. Results

1) *Main Result: Pareto Dominance*: We compare the proposed advantage-regression router against the heuristic routing baselines. Figure 5 shows that a single trained router traces a continuous Pareto frontier that consistently dominates all baselines across all three benchmarks. In particular, the advantage of the proposed routing strategy is most pronounced in the




Fig. 6. Router ranking reliability. Validation images are grouped into deciles based on the predicted advantage $\hat{A}$ (5-seed ensemble). Each point represents the mean true advantage $A = L_{\text{base}} - L_{\text{super}}$ within that decile (error bars denote standard error of the mean). The mean true advantage increases consistently with the predicted decile across all three datasets, demonstrating that a higher router score reliably identifies frames that benefit more from the deep path.




Fig. 7. Per-image advantage $A = L_{\text{base}} - L_{\text{super}}$ distribution on the validation sets. Both are concentrated near zero with a long positive tail—most frames gain little from the deep path while a minority gain a lot. The tail is wider on KITTI ($\sigma=0.21$) than on BDD100K ($\sigma=0.10$).

low-budget regime, where selecting the right subset of frames for deeper execution is critical.

Table III provides the corresponding operating points along the Pareto curve. On KITTI (Figure 5(a)), the router maintains AP@[.50:.95] within 0.1 point of the SUPER-only ceiling (41.4 vs. 41.5) while routing only 67% of frames through the SUPER path. This reduces computation from 26.30 to 23.20 GFLOPs/frame (−11.8%), latency from 20.69 to 18.66 ms (−9.8%), and energy consumption from 3277 to 2957 mJ (−9.8%), while increasing throughput from 48.3 to 53.6 FPS (+11.0%). A more aggressive operating point routes only 43% of frames to the SUPER path, incurring merely a 0.3 AP reduction while lowering computation by 20.6%, latency by 18.6%, and energy by 18.5%.

The same trend holds on BDD100K (Figure 5(b)). At 26% SUPER usage, the router already recovers most of the SUPER-path accuracy gain (22.0 vs. 22.5 AP@[.50:.95]) while re-

ducing computation from 51.72 to 38.02 GFLOPs/frame (−26.5%), latency from 39.16 to 28.79 ms (−26.5%), and energy consumption from 13289 to 9669 mJ (−27.2%). Throughput correspondingly increases from 25.5 to 34.7 FPS (+36.1%).

The same behavior carries over to the large-scale, high-resolution Waymo Open benchmark (Figure 5(c)). At only 35% SUPER usage, the router recovers nearly all of the SUPER-path accuracy gain (38.7 vs. 38.9 AP@[.50:.95], within 0.2 point of the ceiling) while reducing computation from 134.84 to 103.21 GFLOPs/frame (−23.5%), latency from 16.51 to 14.05 ms (−14.9%), and energy consumption from 2108 to 1677 mJ (−20.4%). Throughput correspondingly increases from 60.6 to 71.2 FPS (+17.5%). Notably, the advantage-regression router remains the only strategy that improves monotonically from the very first frames routed to the deep path, whereas the heuristic baselines (most visibly edge density) initially *degrade* accuracy below the always-base anchor by spending the SUPER budget on uninformative frames.

2) Advantage Regression Quality: Because deployment routes by thresholding the predicted advantage $\hat{A}$, the router need not predict the exact advantage magnitude; instead, it must correctly rank frames according to how much they benefit from the deep path. Figure 6 evaluates this property by grouping validation images into deciles based on $\hat{A}$ and reporting the mean true advantage $A = L_{\text{base}} - L_{\text{super}}$ within each group. On all three datasets, the mean true advantage increases consistently with the predicted decile, indicating that frames assigned higher scores indeed benefit more from deeper execution. This confirms that the router successfully captures the relative utility of the deep path, which is the key requirement for threshold-based routing. We emphasize ranking rather than per-image regression accuracy because the single-image target is inherently noisy, making robust ordering a more meaningful indicator of deployment performance.

The pronounced increase in the highest decile is explained




Fig. 8. Runtime latency-budget tracking under realistic scene-condition drift on BDD100K MOT, run as a live on-device closed loop. Sequences are concatenated in alternating condition order (night↔daytime, city↔highway, clear↔rainy; rows), with shaded bands marking each segment’s condition. A PI controller adjusts the router threshold τ to hold a step (left) and a sawtooth (right) latency budget $L^*(t)$ (black dashed); the realized mean latency (red) tracks it with MAE 0.18–0.30 ms across all condition drifts.

by the distribution of the true advantage itself (Figure 7). The advantage is heavily concentrated near zero while exhibiting a long positive tail, indicating that most frames gain little from deeper execution whereas a small subset benefits substantially. The router successfully concentrates these rare high-payoff frames into the top-ranked decile, enabling limited computational resources to be allocated where deeper execution yields the greatest return. This positive tail is noticeably wider on KITTI ($\sigma=0.21$) than on BDD100K ($\sigma=0.10$). Accordingly, the advantage distribution on BDD100K spans a much narrower range, as also reflected by the different y-axis scales in Figure 6, meaning that most achievable operating points are covered by a relatively narrow range of threshold values τ .

3) *Runtime Latency-Budget Tracking*: The single-scalar threshold τ exposes the operating point as a tunable dial, allowing the detector to be steered at runtime to track a time-varying latency budget $L^*(t)$ (e.g., a thermal- or battery-driven schedule) without any retraining. We wrap τ in a proportional-integral (PI) control loop: the controller compares the target budget against an exponential moving average (EMA) of the realized per-frame latency and nudges τ to cancel the error. To stress the controller under realistic non-stationary condi-

tions rather than abrupt synthetic boundaries, we concatenate BDD100K MOT sequences with alternating scene conditions (night↔daytime, city↔highway, clear↔rainy). This forces the input distribution—and hence the expected advantage of the deep path—to drift continuously as the scene evolves.

Because individual frames strictly execute at discrete latencies (the super or base), no single frame can exactly match an intermediate target budget. Thus, the controllable quantity is the mean latency, effectively the fraction of frames routed to the super path, which the PI controller regulates using the EMA error. Figure 8 evaluates this runtime tracking as a live on-device closed loop. To visualize the tracked mean without the artificial phase delay inherently introduced by causal trailing filters, we smooth the binary per-frame measurements using a zero-phase (centered) moving average. This purely offline visualization accurately reflects the controller’s tracking quality without smoothing-induced lag, while the underlying closed loop remains unchanged. The realized mean latency tracks both step and sawtooth target budgets with exceptional precision across all condition drifts, achieving a mean absolute error of only 0.18–0.30 ms. Ultimately, a single trained router successfully turns a static Pareto front into a

TABLE IV

TENSORRT FP16 ENGINE-SWITCHING OVERHEAD ON BDD100K (RTX 3090, PURE GPU INFERENCE, CUDA-EVENT TIMED). ALTERNATING FORCES A BASE $\leftrightarrow$ SUPER ENGINE SWITCH EVERY FRAME (WORST CASE); ITS LATENCY IS WITHIN 1% OF THE MEAN OF THE TWO STATIC SINGLE-PATH ENGINES, SO ROUTING BETWEEN TWO RESIDENT ENGINES IS EFFECTIVELY FREE.

Engine schedule	Latency (ms)	FPS
BASE (always-base)	6.09	164.3
SUPER (always-super)	9.70	103.0
ALTERNATING (switch every frame)	7.97	125.5
mean(base, super)	7.89	—
switch overhead	+0.08 ms	(+1.0%)

robust, runtime-adjustable system that seamlessly adapts to dynamic environments.

4) *TensorRT Deployment*: A natural concern is whether depth routing remains practical on a production inference stack: TensorRT compiles a *static* graph and cannot conditionally skip layers inside a single engine. We therefore realize routing by exporting the BASE and SUPER paths as two separate FP16 engines and keeping both resident in GPU memory; the router simply selects which engine to execute per frame. The only routing-specific concern is then the cost of *switching* between the two resident engines every time the chosen path changes. We measure this on BDD100K (RTX 3090) by comparing each single-path engine run in isolation against an ALTERNATING schedule that forces an engine switch on every frame—a worst case far more frequent than any realistic routing trace (Table IV). The alternating schedule costs only +1.0% over the mean of the two single-path runs, confirming that engine switching is effectively free and that the two-engine design adds no meaningful deployment overhead.

One deployment caveat concerns the operating threshold itself. FP16 quantization slightly perturbs the router’s predicted advantage $\hat{A}$, so a threshold τ calibrated in FP32 realizes a different SUPER usage once the engines run in FP16 (e.g., on BDD100K, $\tau=0.086$ routes 14.8% of frames to the deep path in FP32 but 5.9% in FP16). Crucially, quantization shifts only the *calibration* of τ , not the ranking it induces: $\hat{A}$ remains monotone, so the same accuracy–efficiency frontier is reachable, merely at different threshold values. Recovering a target operating point therefore requires no retraining—only a one-pass recalibration of the τ -to-budget map on a small held-out clip evaluated through the FP16 engines. In the budget-tracking deployment (Sec. IV-B3) even this step is unnecessary: the PI controller adjusts τ online from the realized latency, absorbing the FP16 shift automatically.

C. Ablation Studies

1) *Input Feature: Backbone vs. Neck vs. Both*: We compare three feature sources for the router on both datasets: the backbone pyramid features (input-level representations, before the detection neck), the neck features (prediction-level representations, immediately before the heads), and their concatenation (both).




Fig. 9. Feature-source ablation on BDD100K (AP@[.50:.95] vs. SUPER usage; mean $\pm$ std over 5 seeds). The three sources are close, but the both-feature router attains the best-or-tied mean with the tightest seed band, whereas backbone-only has the widest variance.




Fig. 10. Router design and grid resolution on KITTI. GAP-MLP is compared with TinyConv at 2×2 and 8×8 grids in AP@.50:.95.

As illustrated in Figure 9, the mean accuracy–efficiency trajectories of all three configurations are largely comparable on the BDD100K dataset. However, analyzing the performance variance across multiple training seeds reveals a critical distinction in routing stability. Relying solely on the backbone yields the widest variance band (red shaded area), suggesting that early-stage representations alone lack the task-specific refinement necessary for consistent advantage prediction. Conversely, the concatenation of both feature sets (green curve) consistently attains the best-or-tied mean performance across the entire operational range while exhibiting the tightest variance. By providing the router simultaneously with coarse, low-level spatial context from the backbone and refined, task-specific semantics from the neck, the system achieves the most robust routing behavior. Given that this concatenation introduces negligible computational overhead (as previously established in Table II), we confidently adopt the *both* configuration as our default design.

2) *Router Design and Grid Resolution*: We compare GAP-MLP with TinyConv at 2×2 and 8×8 grids (Figure 10). As

shown in the performance trajectories, the TinyConv architecture with a coarse 2×2 pooling grid consistently achieves the optimal accuracy–efficiency trade-off, outperforming both the spatially-agnostic GAP-MLP and the higher-resolution 8×8 variant. The gap between GAP-MLP and TinyConv 2×2 highlights the importance of spatial awareness. Because GAP-MLP globally collapses the feature maps, it discards all spatial layout information. The superiority of the 2×2 grid indicates that preserving a coarse geometric context—allowing the router to roughly distinguish where dense or difficult objects might be located in the frame—provides valuable cues for predicting the advantage of the deep path. Conversely, increasing the spatial grid resolution to 8×8 proves detrimental. Not only does it yield the lowest mean performance, but it also suffers from severe instability across varying training seeds, as evidenced by its conspicuously wide variance band (green shaded area). This suggests that forcing a lightweight router to process overly fine-grained spatial details causes it to overfit to local feature patterns rather than learning a generalized representation of overall scene difficulty. Consequently, the 2×2 grid emerges as the sweet spot, providing just enough spatial layout to inform routing decisions without introducing the optimization challenges of higher resolutions.

3) *Previous-Action Sampling Ratio $p(\text{super})$* : During training, we sample the assumed previous action (the path whose features the router conditions on) as SUPER with probability p , sweeping $p \in \{0.0, 0.5, 1.0\}$ (Figure 11). We find that the balanced setting $p = 0.5$ yields the best performance, outperforming both extreme configurations. This behavior stems from a train-serve distribution matching effect. At deployment, the router must read features produced by whichever path executed on the previous frame, whether it was BASE or SUPER, and the same image yields slightly different feature maps under each path. If training only exposes SUPER path features ($p = 1.0$), the router fails to calibrate to BASE path features and mis-ranks frames whenever the previous frame ran BASE. Conversely, $p = 0.0$ fails symmetrically. Training with a balanced $p = 0.5$ exposes the router to both feature distributions equally, ensuring that its advantage prediction remains calibrated regardless of the upstream routing decision. The takeaway is that the previous action is an essential component of the router’s input distribution, and balanced sampling is vital for robust, causal per-frame deployment.

V. CONCLUSION

We reformulated dynamic-depth routing as advantage regression, decoupling the training objective from the deployment operating point. A single lightweight router predicts the per-image super advantage from the detector’s intermediate features, and a threshold sweep recovers an entire continuous accuracy–efficiency Pareto curve with no trade-off weight and no anti-collapse term. On KITTI-tracking the router Pareto-dominates random, luminance, edge-density, and confidence baselines and comes within 0.1 AP of super-only accuracy at a 12% FLOPs (and 10% energy) reduction, at negligible overhead. A limitation is that the regression view targets the




Fig. 11. Previous-action sampling ratio $p(\text{super})$ during training (AP@.50:.95) on KITTI.

two-path (super/base) decision; extending it to multi-level or per-layer (2^N) routing, where the decision space is no longer enumerable and sampling- or marginalization-based estimators become necessary, is left to future work. A further limitation is that, with a binary per-frame actuator, what we regulate is the *average* latency (throughput/energy/thermal budgets); hard per-frame real-time deadlines are guaranteed only when the super path itself meets the deadline, and tighter per-frame guarantees would require finer-grained multi-level depths, also left to future work.

REFERENCES

- [1] W. Kang, H. Lee, and J. Lee, “Anydepth-detr/-yolo: Any-depth object detection with a single network,” *arXiv preprint arXiv:2605.09407*, 2026.
- [2] Z. Lin, Y. Wang, J. Zhang, and X. Chu, “Dynamicdet: A unified dynamic architecture for object detection,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2023, pp. 6282–6291.
- [3] D. Kuhse, H. Teper, S. Buschjäger *et al.*, “You only look once at anytime (AnytimeYOLO): Analysis and optimization of early-exits for object-detection,” *arXiv preprint arXiv:2503.17497*, 2025.
- [4] Y. Guo, H. Shi, A. Kumar, K. Grauman, T. Rosing, and R. Feris, “Spot-tune: transfer learning through adaptive fine-tuning,” in *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, 2019, pp. 4805–4814.
- [5] Y. Meng, C.-C. Lin, R. Panda, P. Sattigeri, L. Karlinsky, A. Oliva, K. Saenko, and R. Feris, “Ar-net: Adaptive frame resolution for efficient action recognition,” in *European conference on computer vision*. Springer, 2020, pp. 86–104.
- [6] Z. Wu, T. Nagarajan, A. Kumar, S. Rennie, L. S. Davis, K. Grauman, and R. Feris, “Blockdrop: Dynamic inference paths in residual networks,” in *Proceedings of the IEEE conference on computer vision and pattern recognition*, 2018, pp. 8817–8826.
- [7] S. Teerapittayanon, B. McDanel, and H. Kung, “BranchyNet: Fast inference via early exiting from deep neural networks,” in *Proc. Int. Conf. Pattern Recognit. (ICPR)*, 2016, pp. 2464–2469.
- [8] G. Huang, D. Chen, T. Li, F. Wu, L. van der Maaten, and K. Weinberger, “Multi-scale dense networks for resource efficient image classification,” in *International Conference on Learning Representations*, 2018. [Online]. Available: <https://openreview.net/forum?id=Hk2aImxAb>
- [9] L. Yang, Y. Han, X. Chen, S. Song, J. Dai, and G. Huang, “Resolution adaptive networks for efficient inference,” in *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, 2020, pp. 2369–2378.

- [10] F. Ilhan, K.-H. Chow, S. Hu, T. Huang, S. Tekin, W. Wei, Y. Wu, M. Lee, R. Kompella, H. Latapie *et al.*, "Adaptive deep neural network inference optimization with eenet," in *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, 2024, pp. 1373–1382.
- [11] M. Figurnov, M. D. Collins, Y. Zhu, L. Zhang, J. Huang, D. Vetrov, and R. Salakhutdinov, "Spatially adaptive computation time for residual networks," in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2017, pp. 1039–1048.
- [12] J. Yu, L. Yang, N. Xu, J. Yang, and T. Huang, "Slimmable neural networks," in *International Conference on Learning Representations*, 2019. [Online]. Available: <https://openreview.net/forum?id=H1gMCsAqY7>
- [13] J. Yu and T. S. Huang, "Universally slimmable networks and improved training techniques," in *Proceedings of the IEEE/CVF international conference on computer vision*, 2019, pp. 1803–1811.
- [14] H. Cai, C. Gan, T. Wang, Z. Zhang, and S. Han, "Once for all: Train one network and specialize it for efficient deployment," in *International Conference on Learning Representations*, 2020.
- [15] X. Wang, F. Yu, Z.-Y. Dou, T. Darrell, and J. E. Gonzalez, "Skipnet: Learning dynamic routing in convolutional networks," in *Proceedings of the European conference on computer vision (ECCV)*, 2018, pp. 409–424.
- [16] W. Kang and H. Lee, "Adaptive depth networks with skippable sub-paths," *Advances in Neural Information Processing Systems*, vol. 37, pp. 33 213–33 231, 2024.
- [17] L. Yang, Z. Zheng, J. Wang, S. Song, G. Huang, and F. Li, "AdaDet: An adaptive object detection system based on early-exit neural networks," *IEEE Trans. Cogn. Develop. Syst.*, vol. 16, no. 1, pp. 332–345, 2024.
- [18] S. Heo, S. Cho, Y. Kim, and H. Kim, "Real-time object detection system with multi-path neural networks," in *Proc. IEEE Real-Time Embedded Technol. Appl. Symp. (RTAS)*, 2020, pp. 174–187.
- [19] A. Soyigit, S. Yao, and H. Yun, "VALO: A versatile anytime framework for LiDAR-based object detection deep neural networks," *IEEE Trans. Comput.-Aided Design Integr. Circuits Syst.*, vol. 43, no. 11, pp. 4045–4056, 2024.
- [20] X. Dai, D. Chen, M. Liu, Y. Chen, and L. Yuan, "Da-nas: Data adapted pruning for efficient neural architecture search," in *European conference on computer vision*. Springer, 2020, pp. 584–600.
- [21] N. Shazeer, A. Mirhoseini, K. Maziarz, A. Davis, Q. Le, G. Hinton, and J. Dean, "Outrageously large neural networks: The sparsely-gated mixture-of-experts layer," in *International Conference on Learning Representations*, 2017. [Online]. Available: <https://openreview.net/forum?id=B1ckMDqlg>
- [22] W. Fedus, B. Zoph, and N. Shazeer, "Switch transformers: Scaling to trillion parameter models with simple and efficient sparsity," *Journal of Machine Learning Research*, vol. 23, no. 120, pp. 1–39, 2022.
- [23] C. Riquelme, J. Puigcerver, B. Mustafa, M. Neumann, R. Jenatton, A. Susano Pinto, D. Keysers, and N. Houlsby, "Scaling vision with sparse mixture of experts," *Advances in Neural Information Processing Systems*, vol. 34, pp. 8583–8595, 2021.
- [24] A. Geiger, P. Lenz, and R. Urtasun, "Are we ready for autonomous driving? the kitti vision benchmark suite," in *2012 IEEE conference on computer vision and pattern recognition*. IEEE, 2012, pp. 3354–3361.
- [25] F. Yu, H. Chen, X. Wang, W. Xian, Y. Chen, F. Liu, V. Madhavan, and T. Darrell, "Bdd100k: A diverse driving dataset for heterogeneous multitask learning," in *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, 2020, pp. 2636–2645.
- [26] P. Sun, H. Kretzschmar, X. Dotiwalla, A. Chouard, V. Patnaik, P. Tsui, J. Guo, Y. Zhou, Y. Chai, B. Caine *et al.*, "Scalability in perception for autonomous driving: Waymo open dataset," in *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, 2020, pp. 2446–2454.